

Preprocess.py

1. Detect_keypoints():
 - a. Given the image file path, read the image.
 - b. Using cv2.SIFT, get the keypoints of the image.
2. Create_feature_matches():
 - a. Using cv2.BFMatcher, feature match descriptors from each image using k_nearest_neighbour value of 2.
 - b. Filter the feature matches using the Lowe ratio test using a ratio of 0.6.
3. Create_ransac_matches():
 - a. Using cv2.findEssentialMatrix, calculate the essential matrix and the inlier mask to differentiate inliers and outliers.
4. Create_scene_graph():
 - a. For each image pair, if the number of inlier feature matches exceeds the minimum number of inliers threshold, we add an edge between those 2 images' nodes.

Sfm.py

1. Get_init_image_ids():
 - a. For each pair of images, select the image pair that has the highest number of feature matches.
2. Get_init_extrinsics():
 - a. Using cv2.recoverPose, recover the rotation and translation vectors from essential matrix.
 - b. Using cv2.rodrigues, convert rotation vector to the rotation matrix.
 - c. Create and return the extrinsic matrix by appending translation vector to rotation matrix $[R|t]$.
3. Get_next_pair():
 - a. For each image id in registered_ids, check each of its adjacent nodes (which is not registered) and select and return the pair with the highest number of feature match inliers.
4. Solve_pnp():
 - a. Using cv2.solvePnP, calculate the rotation and translation vectors from the camera intrinsic and selected 3d/2d points.
 - b. Using cv2.rodrigues, convert rotation vector to the rotation matrix.
 - c. Use get_reprojection_residuals() to calculate the reprojection residuals
 - i) Get_reprojection_residuals():
 - (1) Create the homography using intrinsic and extrinsic matrices.
 - (2) For each point in points3d, reproject it onto the 2d plane using homography matrix.
 - (3) Calculate the Euclidean distance between reprojected 2d point and original 2d point.
 - (4) Return the residuals once all points3d have been reprojected.
5. Add_points3d():
 - a. Given the unregistered feature matches between 2 images, triangulate these corresponding points to create new 3d points.

Bundle_adjustment.py

1. Compute_ba_residuals():
 - a. Calculate the different homography matrices from the intrinsic and extrinsic matrices.
 - b. Homogenise the 3d points given by appending a $[1, \dots]^T$ vector to the 3d points.
 - c. Retrieve the corresponding 3d points using points3d[points3d_idxs].
 - d. Reproject the corresponding 3d points using the homography matrices calculated before and normalise the new reprojected 2d points.
 - e. Calculate the difference between the (x-axis values of the reprojected points and initial 2d points) squared.

- f. Calculate the difference between the (y-axis values of the reprojected points and initial 2d points) squared.
- g. Square root the x-axis/y-axis differences to get the Euclidean distances between initial and reprojected 2d points across the different camera poses.
- h. Sum the Euclidean distances across the different camera poses to get the ba residuals.

Results




Figure 1 Mini-temple w/o BA



Figure 1 TA's Mini-temple w/o BA



Figure 3 Mini-temple with BA



Figure 4 TA's Mini-temple with BA




Figure 5 Temple w/o BA




Figure 6 TA's Temple w/o BA